

# Project 3: Semantic Segmentation

## Part 1:

### 1) List of configs used:

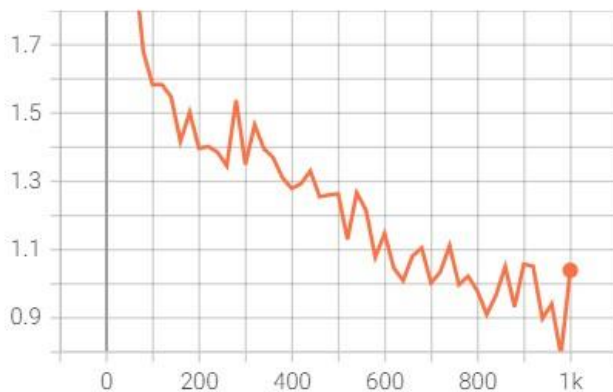
For my model I used the “`faster_rcnn_R_101_FPN_3x.yaml`” as suggested, but I changed some configurations changes for high accuracy

Firstly, I increased MAX\_ITER to 1000 and then reduced BATCH\_SIZE to 32. I let rest of the configurations to remain same.


2) Just by increasing the number of iterations, the accuracy increased as the model was trained longer. Reducing batchsize per image also helped a lot as a smaller number of images ran through the model at a time, hence they were generalized less which also reduced loss while training longer.

3) Plotting total training loss and accuracy for the model

total\_loss  
tag: total\_loss



fast\_rcnn/cls\_accuracy  
tag: fast\_rcnn/cls\_accuracy



### 4) Visualizing test sample:



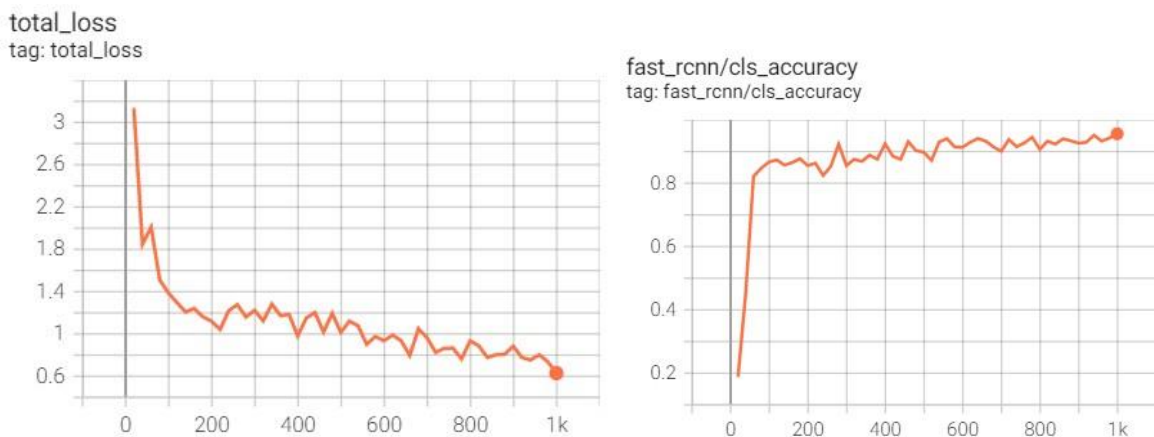




From the test images we can note that while the model is able to object planes in the images, there are also misidentifications along with failing to detect some images made by the model. As seen in the image 2, the model misidentified 2 corners of a helicopter pad as planes. Whereas in image 3, the model failed to detect some planes.

- 5) Now conducting an ablation study trying to understand how batchsize affect a model. We can run our model again with BATCHSIZE set as 512 and see what happens.

After training the new model, we can get plot its loss and accuracy as



Comparing the total\_loss and accuracy plots for the models, we see that batchsize512 model's plots for total\_loss and accuracy are very smooth and gradual compared to batchsize32 model. However, looking at the plots we can also say the accuracy of batchsize32 model is a bit better.

If we look the APs for both the models

```
[11/07 08:12:55 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[11/07 08:12:55 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.18 seconds.
[11/07 08:12:55 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[11/07 08:12:55 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.240
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.429
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.243
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.149
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.305
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.551
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.114
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.270
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.155
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.336
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.676
[11/07 08:12:55 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
  AP    AP50   AP75   APs    APm    AP1
-----:-----:-----:-----:-----:-----:
24.017 | 42.879 | 24.266 | 14.916 | 30.493 | 55.077

[11/07 07:20:58 d2.evaluation.fast_eval_api]: Evaluate annotation type "bbox"
[11/07 07:20:58 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.20 seconds.
[11/07 07:20:58 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[11/07 07:20:58 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.238
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.257
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.137
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.559
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.014
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.115
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.267
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.136
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.347
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.693
[11/07 07:20:58 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
  AP    AP50   AP75   APs    APm    AP1
-----:-----:-----:-----:-----:-----:
23.832 | 40.276 | 25.665 | 13.725 | 31.794 | 55.903
```

AP50 of BATCHSIZE32 model is much better than the other model.

Getting the same test images for BATCHSIZE512 model and comparing them with BATCHSIZE32







From the test images it becomes clear that both the models perform almost similarly and made the same mistakes. However, prediction accuracies are slightly better for BATCH32 model.

## Part 2:

- 1) For training the model, I set up hyperparameters as `num_epochs = 5`, `batch_size = 64`, `learning_rate = 0.001`, `weight_decay = 1e-5` and I used Adam optimizing algorithm.
- 2) In the network, first I add more down layers going from 3 channels to 16 and then 32, 64, 128, 256, 512 and 1024. Then, to match with the downsampling I added up layers from 1024 back to 3. I also added skip connections between corresponding channels. Lastly, I created a new class called block which pretty much does (convolution, batchnorm, relu)x2 for same in/out channels just to have more convolution the networks.

Vaibhav Saini  
301386847

Late days used: 2  
Team Name: 301386847


Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 128, 128]	448
BatchNorm2d-2	[-1, 16, 128, 128]	32
ReLU-3	[-1, 16, 128, 128]	0
conv-4	[-1, 16, 128, 128]	0
Conv2d-5	[-1, 16, 128, 128]	2,320
BatchNorm2d-6	[-1, 16, 128, 128]	32
ReLU-7	[-1, 16, 128, 128]	0
Conv2d-8	[-1, 16, 128, 128]	2,320
BatchNorm2d-9	[-1, 16, 128, 128]	32
ReLU-10	[-1, 16, 128, 128]	0
block-11	[-1, 16, 128, 128]	0
Conv2d-12	[-1, 32, 128, 128]	4,640
BatchNorm2d-13	[-1, 32, 128, 128]	64
ReLU-14	[-1, 32, 128, 128]	0
conv-15	[-1, 32, 128, 128]	0
MaxPool2d-16	[-1, 32, 64, 64]	0
down-17	[-1, 32, 64, 64]	0
Conv2d-18	[-1, 32, 64, 64]	9,248
BatchNorm2d-19	[-1, 32, 64, 64]	64
ReLU-20	[-1, 32, 64, 64]	0
Conv2d-21	[-1, 32, 64, 64]	9,248
BatchNorm2d-22	[-1, 32, 64, 64]	64
ReLU-23	[-1, 32, 64, 64]	0
block-24	[-1, 32, 64, 64]	0
Conv2d-25	[-1, 64, 64, 64]	18,496
BatchNorm2d-26	[-1, 64, 64, 64]	128
ReLU-27	[-1, 64, 64, 64]	0
conv-28	[-1, 64, 64, 64]	0
MaxPool2d-29	[-1, 64, 32, 32]	0
down-30	[-1, 64, 32, 32]	0
Conv2d-31	[-1, 64, 32, 32]	36,928
BatchNorm2d-32	[-1, 64, 32, 32]	128
ReLU-33	[-1, 64, 32, 32]	0
Conv2d-34	[-1, 64, 32, 32]	36,928
BatchNorm2d-35	[-1, 64, 32, 32]	128
ReLU-36	[-1, 64, 32, 32]	0
block-37	[-1, 64, 32, 32]	0
Conv2d-38	[-1, 128, 32, 32]	73,856
BatchNorm2d-39	[-1, 128, 32, 32]	256
ReLU-40	[-1, 128, 32, 32]	0
conv-41	[-1, 128, 32, 32]	0
MaxPool2d-42	[-1, 128, 16, 16]	0
down-43	[-1, 128, 16, 16]	0
Conv2d-44	[-1, 128, 16, 16]	147,584
BatchNorm2d-45	[-1, 128, 16, 16]	256
ReLU-46	[-1, 128, 16, 16]	0
Conv2d-47	[-1, 128, 16, 16]	147,584
BatchNorm2d-48	[-1, 128, 16, 16]	256
ReLU-49	[-1, 128, 16, 16]	0
block-50	[-1, 128, 16, 16]	0
Conv2d-51	[-1, 256, 16, 16]	295,168
BatchNorm2d-52	[-1, 256, 16, 16]	512
ReLU-53	[-1, 256, 16, 16]	0
conv-54	[-1, 256, 16, 16]	0
MaxPool2d-55	[-1, 256, 8, 8]	0
down-56	[-1, 256, 8, 8]	0
Conv2d-57	[-1, 256, 8, 8]	590,080
BatchNorm2d-58	[-1, 256, 8, 8]	512
ReLU-59	[-1, 256, 8, 8]	0
Conv2d-60	[-1, 256, 8, 8]	590,080
BatchNorm2d-61	[-1, 256, 8, 8]	512
ReLU-62	[-1, 256, 8, 8]	0
block-63	[-1, 256, 8, 8]	0
Conv2d-64	[-1, 512, 8, 8]	1,180,160
BatchNorm2d-65	[-1, 512, 8, 8]	1,024
ReLU-66	[-1, 512, 8, 8]	0
conv-67	[-1, 512, 8, 8]	0
MaxPool2d-68	[-1, 512, 4, 4]	0
down-69	[-1, 512, 4, 4]	0
Conv2d-70	[-1, 512, 4, 4]	2,359,808
BatchNorm2d-71	[-1, 512, 4, 4]	1,024
ReLU-72	[-1, 512, 4, 4]	0
Conv2d-73	[-1, 512, 4, 4]	2,359,808
BatchNorm2d-74	[-1, 512, 4, 4]	1,024
ReLU-75	[-1, 512, 4, 4]	0
block-76	[-1, 512, 4, 4]	0
Conv2d-77	[-1, 1024, 4, 4]	4,719,616
BatchNorm2d-78	[-1, 1024, 4, 4]	2,048
ReLU-79	[-1, 1024, 4, 4]	0
conv-80	[-1, 1024, 4, 4]	0
MaxPool2d-81	[-1, 1024, 2, 2]	0
down-82	[-1, 1024, 2, 2]	0
ConvTranspose2d-83	[-1, 1024, 4, 4]	4,195,328
Conv2d-84	[-1, 512, 4, 4]	4,719,104
BatchNorm2d-85	[-1, 512, 4, 4]	1,024
ReLU-86	[-1, 512, 4, 4]	0
conv-87	[-1, 512, 4, 4]	0
up-88	[-1, 512, 4, 4]	0
BatchNorm2d-89	[-1, 512, 4, 4]	1,024
ConvTranspose2d-90	[-1, 512, 8, 8]	1,049,088
Conv2d-91	[-1, 256, 8, 8]	1,179,904
BatchNorm2d-92	[-1, 256, 8, 8]	512
ReLU-93	[-1, 256, 8, 8]	0
conv-94	[-1, 256, 8, 8]	0
up-95	[-1, 256, 8, 8]	0
BatchNorm2d-96	[-1, 256, 8, 8]	512
ConvTranspose2d-97	[-1, 256, 16, 16]	262,400
Conv2d-98	[-1, 128, 16, 16]	295,040
BatchNorm2d-99	[-1, 128, 16, 16]	256
ReLU-100	[-1, 128, 16, 16]	0
conv-101	[-1, 128, 16, 16]	0
up-102	[-1, 128, 16, 16]	0
BatchNorm2d-103	[-1, 128, 16, 16]	256
ConvTranspose2d-104	[-1, 128, 32, 32]	65,664
Conv2d-105	[-1, 64, 32, 32]	73,792
BatchNorm2d-106	[-1, 64, 32, 32]	128
ReLU-107	[-1, 64, 32, 32]	0
conv-108	[-1, 64, 32, 32]	0
up-109	[-1, 64, 32, 32]	0
BatchNorm2d-110	[-1, 64, 32, 32]	128
ConvTranspose2d-111	[-1, 64, 64, 64]	16,448
Conv2d-112	[-1, 32, 64, 64]	18,464
BatchNorm2d-113	[-1, 32, 64, 64]	64
ReLU-114	[-1, 32, 64, 64]	0
conv-115	[-1, 32, 64, 64]	0
up-116	[-1, 32, 64, 64]	0
BatchNorm2d-117	[-1, 32, 64, 64]	64
ConvTranspose2d-118	[-1, 32, 128, 128]	4,128
Conv2d-119	[-1, 16, 128, 128]	4,624
BatchNorm2d-120	[-1, 16, 128, 128]	32
ReLU-121	[-1, 16, 128, 128]	0
conv-122	[-1, 16, 128, 128]	0
up-123	[-1, 16, 128, 128]	0
BatchNorm2d-124	[-1, 16, 128, 128]	32
Conv2d-125	[-1, 3, 128, 128]	435
BatchNorm2d-126	[-1, 3, 128, 128]	6
ReLU-127	[-1, 3, 128, 128]	0
conv-128	[-1, 3, 128, 128]	0
BatchNorm2d-129	[-1, 3, 128, 128]	6
Conv2d-130	[-1, 1, 128, 128]	28
conv-131	[-1, 1, 128, 128]	0

3) I kept the loss function same as baseline using BCEWithLogitsLoss. Plotting epochs vs total training loss.

4) The final IoU for the model is:

#images: 7980, Mean IoU: 0.7702353964455629

5) Testing Samples:






Part 3:

- 1) Kaggle teamname: 301386847
- 2) Score on Kaggle leaderboard: 0.46282
- 3) Test Samples

Vaibhav Saini  
301386847




Late days used: 2  
Team Name: 301386847





Vaibhav Saini  
301386847

Late days used: 2  
Team Name: 301386847



## Part 4:

### 1) Visualizations

#### Segs evaluation

```
[11/05 10:21:35 d2.evaluation.fast_eval_api]: Evaluate annotation type *segm*
[11/05 10:21:36 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.89 seconds.
[11/05 10:21:36 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[11/05 10:21:36 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.077
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.272
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.012
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.036
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.087
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.317
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.007
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.052
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.106
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.050
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.129
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.374
[11/05 10:21:36 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP | AP50 | AP75 | APs | APm | APl |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| 7.669 | 27.247 | 1.248 | 3.618 | 8.737 | 31.681 |
```


#### Bbox evaluation

```
[11/05 10:21:34 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[11/05 10:21:35 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate() finished in 0.23 seconds.
[11/05 10:21:35 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[11/05 10:21:35 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.01 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.278
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.483
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.290
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.173
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.360
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.581
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.015
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.123
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.313
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.180
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.398
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.711
[11/05 10:21:35 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| 27.781 | 48.258 | 28.988 | 17.347 | 35.970 | 58.070 |
```









- 2) Comparing this mask RCNN model with Part3 model, we can see that our part3 model was actually able to perform better than mask RCNN model. With the masks making fewer mistakes and misclassifications.

## References:

[My experiment with UNet - building an image segmentation model \(analyticsindiamag.com\)](https://analyticsindiamag.com)

[Gentle Introduction to the Adam Optimization Algorithm for Deep Learning \(machinelearningmastery.com\)](https://machinelearningmastery.com)

[python - What does "unsqueeze" do in Pytorch? - Stack Overflow](https://stackoverflow.com)

[Detectron2 Train a Instance Segmentation Model \(gilberttanner.com\)](https://gilberttanner.com)

[Quick intro to Instance segmentation: Mask R-CNN \(kharshit.github.io\)](https://kharshit.github.io)

[Use Custom Datasets — detectron2 0.6 documentation](https://detectron2.readthedocs.io)

[detectron2/MODEL\\_ZOO.md at main · facebookresearch/detectron2 · GitHub](https://github.com/facebookresearch/detectron2)

[Object Detection for Dummies Part 1: Gradient Vector, HOG, and SS \(lilianweng.github.io\)](https://lilianweng.github.io)

[Understanding Semantic Segmentation with UNET | by Harshall Lamba | Towards Data Science](#)

[Metrics to Evaluate your Semantic Segmentation Model | by Ekin Tiu | Towards Data Science](#)

[python - Google Colab is very slow compared to my PC - Stack Overflow](https://stackoverflow.com)